

Replit Agent Prompt — Fix Blog Post Indexability, Canonical & Metadata

Context

You are working on the Job Genie web app (React, hosted on Replit; production domain `https://job-genie.ai`). The app has a blog at `/blog/{slug}`. Blog posts are failing three critical SEO/AEO requirements that make them invisible to search engines and AI answer-engine crawlers (most of which do not execute JavaScript). Fix all three without breaking the homepage or any existing route.

Reproduction: Fetching `/blog/stay-motivated-applying-jobs-months-unemployment` returns only the `<head>` metadata — the article body is absent from the server response, the `<link rel="canonical">` points to the homepage (`https://job-genie.ai/`), and the `<title>` / meta description / OG / Twitter tags are the homepage's values instead of the post's.

Objectives

1. The full article body must be present in the **initial server response** (raw HTML), not injected client-side after hydration.
2. Each blog post must emit a **self-referential canonical**: `https://job-genie.ai/blog/{slug}`.
3. Each blog post must emit **its own** `<title>`, meta description, OG, and Twitter tags from that post's data — never the homepage defaults.

Step 1 — Diagnose before changing anything

Report back before editing:

- The framework and rendering mode (Next.js? Vite + React SPA? Express + client React? Other?).
- How blog posts are stored (markdown + YAML front matter? CMS? database?) and where those files/records live.
- The exact component/function that generates `<head>` tags, and **why** it currently outputs homepage values on blog routes.
- Whether blog routes are statically generated, server-rendered, or client-only.

Do not proceed to edits until these are identified.

Step 2 — Fix rendering (body must be in raw HTML)

Choose the approach based on the stack found in Step 1:

- **If Next.js:** render blog posts via SSG (`generateStaticParams` / `getStaticProps`) or SSR so the body ships in the initial HTML.
- **If Vite / CRA SPA (client-only):** add build-time prerendering for all `/blog/*` routes (e.g. a Puppeteer- or `react-snap`-style prerender step), or migrate blog routes to a server-rendered/static path, so each post URL returns complete HTML. Do **not** rely on client hydration for content visibility.

Success target: a `curl` of any post URL returns the full article text.

Step 3 — Fix canonical + per-post metadata

- Replace the global/homepage head logic on blog routes with **per-post** values sourced from each post's front matter/record.
- **Canonical:** `https://job-genie.ai/blog/{slug}` (self-referential). Never the homepage. Use the production domain even when rendered on the `*.replit.dev` preview.
- **Per-post tags:** `<title>`, meta description, `og:title`, `og:description`, `og:url`, `twitter:title`, `twitter:description` all pull from the post. `og:url` = the post's own production URL.
- If a post lacks a custom OG image, falling back to the homepage image is acceptable — but canonical and `og:url` must still be the post's own URL.

Step 4 — Correct existing post data

The `stay-motivated...` post's front matter currently holds the homepage canonical, title, and description (likely copied from the homepage template). Audit **every** blog post's front matter/record and correct:

- `canonical` → its own URL
- `title` → unique, query-matching, post-specific
- `meta-description` → answer-first, post-specific, ~150–160 characters

Step 5 — Add Article schema

Emit `Article` (or `BlogPosting`) JSON-LD on each post: `headline`, `description`, `author`, `datePublished`, `mainEntityOfPage` = post URL. If a post is structured as Q&A, also emit `FAQPage` JSON-LD. Confirm the JSON-LD parses without errors.

Constraints

- Do not alter homepage metadata or break any existing route.
- Do not introduce client-side-only meta injection for content crawlers need to see.
- Keep changes minimal and reversible; list every file you touched with a one-line reason.

Acceptance criteria (run these and show the output)

Run against the preview URL for `/blog/stay-motivated-applying-jobs-months-unemployment`:

1. `curl -s <url> | grep -i 'rel="canonical"'` → shows the post's own URL (`https://job-genie.ai/blog/stay-motivated-applying-jobs-months-unemployment`), **not** the homepage.
2. `curl -s <url> | grep -o '<title>[^<]*'` → shows the post-specific title, not the homepage title.
3. `curl -s <url>` raw HTML contains a verbatim sentence from the article body.
4. An `Article`/`BlogPosting` JSON-LD block is present and valid.

Repeat checks 1–3 on one other blog post to confirm the fix generalizes rather than being hardcoded to one slug.